

TL;DR

Pareton gives AI companies a repeatable way to find and optimize their inference engine — lower GPU-hours, higher throughput, or better latency — without staffing a dedicated inference-optimization team.

We are looking for a small number of design partners to run 2 - 4 week pilots against real workload, hardware, and serving-stack configurations.

Why Now

- Inference demand is growing faster than serving efficiency. At the same time, model architectures, GPU types, request patterns, and serving stacks are changing quickly.
- One-off hand tuning decays. Providers need a repeatable process for discovering, validating, and maintaining workload-specific inference optimizations.
- For large GPU fleets, even 2–5% better serving efficiency or utilization can create meaningful gross-margin impact.

Ideal Design Partner Profiles

1. Inference Providers / Model-Serving Platforms

Examples: Chutes, Together-style platforms, Fireworks-like platforms

Why they care:

- Inference cost directly impacts gross margin.
- They operate vLLM-like stacks or custom forks.
- They constantly benchmark model, hardware, and config changes.
- They want better performance without maintaining a dedicated optimization team.

Best intro targets:

- Head of Infrastructure
- Inference / serving-stack owner

2. NeoClouds / GPU Infrastructure Providers

Examples: Crusoe, Nebius, CoreWeave, Lando

Why they care:

- They operate large GPU fleets where low-single-digit efficiency gains are economically meaningful.
- Inference margin is becoming a competitive lever.
- They may be able to distribute optimization downstream to enterprise customers.

3. Application-Layer AI Companies

Examples: coding assistants, agent companies, enterprise copilots

Why they care:

- Their workloads are often narrower and more stable than generic inference traffic.
- Workload-specific serving can outperform general-purpose inference.
- Cost and latency directly impact product margin and user experience.

Who This Is Especially Relevant For

Companies that:

- Spend materially on LLM inference
- Run or customize vLLM / SGLang / similar serving stacks
- Care about GPU utilization, throughput, latency, or cost per request
- Have repeatable workload patterns
- Do not want to maintain a permanent internal inference-optimization team
- Are open to a short benchmark / pilot using representative workloads

Not A Fit If

- Inference volume is too small for GPU efficiency to matter
- The company only uses fully managed APIs and has no control over serving stack or provider choice
- The workload changes too frequently to define a useful benchmark
- The team cannot share representative workload / hardware / config information, even under NDA

Pareton's Solution

Pareton takes a customer's model, workload profile, hardware, and serving configuration, then runs a structured optimization process to discover improvements that survive benchmarking across realistic deployment conditions.

The current implementation focuses on auditable vLLM code patches and benchmarked serving-stack improvements. Gains that are fragile or environment-specific are filtered out; improvements that hold across agreed test conditions are promoted into the design partner's optimization baseline.

Proof / Early Signals

In an earlier controlled benchmark, the team demonstrated ~35% throughput improvement over a vLLM baseline on a specific production-style workload.

This does not guarantee the same gain in every deployment, but it validates the core thesis: targeted incentivized optimization can improve serving efficiency under defined conditions.

The next step is to run the process against each partner's actual environment and measure real-world savings.

Pilot Proposal

Duration: 2–4 weeks

Design partner provides:

- Representative workload profile
- Target model(s)
- Hardware / GPU type
- Current serving stack and benchmark baseline
- Deployment recommendation
- Priority metric: cost, throughput, latency, utilization, or cost per request

Pareton provides:

- Optimization run against the agreed profile
- Benchmark comparison vs. baseline
- Before/After report
- Commercial proposal if gains are meaningful

Success condition: measurable improvement in throughput, GPU-hour efficiency, latency, utilization, or cost per request under agreed test conditions.

Commercial Model

Pilot → Measurement → Contract

Potential pricing structures:

- Fixed pilot fee + success fee
- Savings-share based on measured efficiency gains
- Subscription for continuous optimization and maintenance

Suggested Intro Framing:

“Pareton is working on workload-specific inference optimization. They are looking for a small number of design partners to test whether their optimization process can reduce GPU-hours or improve throughput for real serving workloads. Worth a technical conversation?”

Contact

Xavier Lyu

xavier@pareton.ai